

▼ AI4FutureWorkForce Data Processing

Exploration of predicting course completion using sklearn and Tensorflow

▼ Dependencies

These are all the libraries we need to run this notebook

```
1 import pandas as pd
2 import plotly.offline as plt
3 import _pickle as pickle
4 import os
5 import numpy as np
6
7 import tensorflow as tf
8
9 import tensorflow.keras as keras
10
11 import plotly.graph_objs as go
12 import itertools
13
14 from sklearn.naive_bayes import GaussianNB
15 from sklearn import svm
16 from sklearn.metrics import accuracy_score
17 from sklearn import linear_model
18 from sklearn.ensemble import RandomForestClassifier
19 from sklearn.model_selection import GridSearchCV
20 from sklearn import preprocessing
21 from sklearn.model_selection import train_test_split
22 from collections import defaultdict
23 from sklearn.metrics import classification_report
24 from numpy.random import seed
25 seed(42)
26 from tensorflow import set_random_seed
27 set_random_seed(42)
28 from tensorflow.python.client import device_lib
29
30 plt.init_notebook_mode(connected=True)
```



```
1 # Verify that the TF device we're using is what we expect
2 device_lib.list_local_devices()
```



```
[name: "/device:CPU:0"
 device_type: "CPU"
 memory_limit: 268435456
 locality {
 }
 incarnation: 5457231463178488201, name: "/device:XLA_CPU:0"
 device_type: "XLA_CPU"
 memory_limit: 17179869184
 locality {
 }
 incarnation: 7095777417938979866
 physical_device_desc: "device: XLA_CPU device"]
```

▼ Load Data

```

1 # Whether or not to use raw or processed dfs
2 PREPROCESS = 0
3 # Whether to fill fee columns or just encode as NaN
4 FILL = 0
5 SHUFFLE = True
6
7 NUM_CLASSES = 2

```

```

1 # List all columns of the tables in data
2 column_list = ['Age (Birthday Masked)', 'Income', 'Education',
3               'MAX(Learner Test Score)', 'Primary Interest In Course',
4               'Hours Coded', 'How Many Hours A Week Can You Commit To Class',
5               'Promise Zone Indicator', 'Hacker Rank Score', 'Location; number']
6
7 # List desired columns for train/test/validation
8 desired_columns = ['Age (Birthday Masked)', 'Income', 'Education',
9                   'MAX(Learner Test Score)', 'Primary Interest In Course',
10                  'Hours Coded', 'How Many Hours A Week Can You Commit To Class',
11                  'Promise Zone Indicator', 'Hacker Rank Score', 'Location; number']
12
13 # Load raw data and run pre-processing
14 if PREPROCESS == 1:
15     df_data = pickle.load( open( "df_data.p", "rb" ) )
16     df_labels = pickle.load( open( "df_labels.p", "rb" ) )
17
18     df_validata = pickle.load( open( "df_validata.p", "rb" ) )
19     df_valilabels = pickle.load( open( "df_valilabels.p", "rb" ) )
20
21     df_valilabels = df_valilabels.add(1)
22     df_labels = df_labels.add(1)
23
24     X_train, X_test, X_val, y_train, y_test, y_val = data_prepare.prepare
25     (column_list, desired_columns, df_data, df_validata, df_labels,
26      df_valilabels, SHUFFLE, FILL)
27
28     print(X_train.shape, X_test.shape, X_val.shape)
29
30 # Otherwise load pre-processed data
31 else:
32     X_train = pickle.load( open( "X_train.p", "rb" ) )
33     X_test = pickle.load( open( "X_test.p", "rb" ) )
34     X_val = pickle.load( open( "X_val.p", "rb" ) )
35
36     y_train = pickle.load( open( "y_train.p", "rb" ) )
37     y_test = pickle.load( open( "y_test.p", "rb" ) )
38     y_val = pickle.load( open( "y_val.p", "rb" ) )

```

Checking that the dataframe looks like it should:

```
1 X_train.head().
```



	Age (Birthday Masked)	Income	MAX(Learner Test Score)	Hours Coded	How Many Hours A Week Can You Commit To Class	Promise Zone Indicator	Location; number
1341	0.353798	1.832671	0.988618	-0.055229	-2.027239	-0.405866	0.571825
1649	-1.263828	-0.299421	0.018359	0.959409	0.381102	-0.405866	0.571825
998	-0.940303	-1.152258	1.312038	0.959409	0.381102	-0.405866	0.571825

```
1 y_train
```

```
↳ array([1, 1, 1, ..., 1, 0, 1])
```

▾ Classification Functions

```
1 class EpochTrack(keras.callbacks.Callback):
2     def on_epoch_end(self, epoch, logs):
3         if epoch % 100 == 0:
4             print('')
5             print('.', end='')
6
7
8 def eval_model(model, test_data, test_labels, VERBOSE, NUM_CLASSES):
9     test_labels_sparse = keras.utils.to_categorical(test_labels, NUM_CLASSES)
10    [loss, mae, acc] = model.evaluate(test_data, test_labels_sparse,
11                                     verbose=VERBOSE)
12
13    print("Mean Abs Error:\t{:7.2f}".format(mae * 1000))
14    print("Loss:\t\t", loss)
15    print("Accuracy:\t", acc)
16
17    preds = model.predict(test_data)
18    y_preds = preds.argmax(axis=-1)
19    print(classification_report(test_labels, y_preds))
20
21    return acc, preds
22
23
24 def plot_history(history):
25     # Define each data series
26     trace1 = go.Scatter(x=history.epoch, y=np.array
27                         (history.history['mean_absolute_error']),
28                         name='Training Loss')
29     trace2 = go.Scatter(x=history.epoch, y=np.array
30                         (history.history['val_mean_absolute_error']),
31                         name='Val Loss')
32
33     # Add each series
34     data = [trace1, trace2]
35
36     # Define graph layout
37     layout = go.Layout(
38         title='Training History',
39         xaxis=dict(title='Epoch'),
40         yaxis=dict(title='Mean Abs Error'))
41
42     fig = go.Figure(data=data, layout=layout)
43     plt.iplot(fig).
44
45 def show_incorrect(y_classes, data, act):
46     incorrects = np.nonzero(y_classes.reshape((-1,)) != act)
47     test_X = data.iloc[incorrects]
48
49     test_y = pd.DataFrame(y_classes[tuple(incorrects)])
50     test_y.columns = ['Prediction']
51
52     act_y = pd.DataFrame(act[tuple(incorrects)])
53     act_y.columns = ['Actual']
54
55     test_X.index = range(len(test_X))
56
57     frames = [test_X, test_y, act_y]
58     df_pred = pd.concat(frames, axis=1)
59
60     return df_pred
61
62
63 def skl_evaluate(model, desired_columns, X, y):
```

```

64     score = model.score(X, y.astype(int))
65
66     preds = model.predict(X)
67     print(classification_report(y, preds))
68
69     print("Accuracy on dataset:\t %f\n" % score)
70
71     importances = model.feature_importances_
72
73     indices = np.argsort(importances)[::-1]
74
75     # Print the feature ranking
76     print("Feature ranking:")
77     for f in range(X.shape[1]):
78         print("%d. %s \t(%f)" % (f + 1, desired_columns[indices[f]], importances[indices[f]]))
79
80     return score

```

▼ Classify

```

1 nb = GaussianNB()
2 nb.fit(X_train, y_train.astype(int))

```

➤ GaussianNB(priors=None, var_smoothing=1e-09)

```

1 X_test.iloc[0, :]

```

➤

Age (Birthday Masked)	-1.048144
Income	0.553416
MAX(Learner Test Score)	0.341779
Hours Coded	-0.055229
How Many Hours A Week Can You Commit To Class	1.585273
Promise Zone Indicator	-0.405866
Location; number	0.571825
Name: 522, dtype: float64	

```

1 nb_score = nb.score(X_test, y_test.astype(int))
2 print("Naive Bayes accuracy on test:\t %f" % nb_score)
3
4 nb_score_val = nb.score(X_val, y_val.astype(int))

```

➤ Naive Bayes accuracy on test: 0.568765

▼ SVM

```

1 clf = svm.SVC(kernel='linear', C=1, random_state=42)
2 clf.fit(X_train, y_train.astype(int))

```

➤ SVC(C=1, cache_size=200, class_weight=None, coef0=0.0, decision_function_shape='ovr', degree=3, gamma='auto_deprecated', kernel='linear', max_iter=-1, probability=False, random_state=42, shrinking=True, tol=0.001, verbose=False)

```

1 clf_score = clf.score(X_test, y_test.astype(int))
2 print("SVM accuracy on test:\t %f" % clf_score)
3

```

```
4 clf_score_val = clf.score(X_val, y_val.astype(int))
```

➞ SVM accuracy on test: 0.596737

▼ Linear Regression

```
1 linr = linear_model.LinearRegression(n_jobs = -1)
2 linr.fit(X_train, y_train.astype(int))
```

➞ LinearRegression(copy_X=True, fit_intercept=True, n_jobs=-1, normalize=False)

```
1 linr_score = linr.score(X_test, y_test.astype(int))
2 print("Linear Regression coefficient of determination, 1 is perfect prediction:\t %f"
3
4 linr_score_val = linr.score(X_val, y_val.astype(int)).
```

➞ Linear Regression coefficient of determination, 1 is perfect prediction: 0.

▼ Random Forest

```
1 rf = RandomForestClassifier(criterion='gini', max_depth=5,
2                             min_samples_leaf=5, min_samples_split=2,
3                             n_estimators = 220, oob_score=True,
4                             max_features=0.5, n_jobs = -1, random_state=42)
5
6 rf.fit(X_train, y_train.astype(int))
```

➞ RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini', max_depth=5, max_features=0.5, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=5, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=220, n_jobs=-1, oob_score=True, random_state=42, verbose=0, warm_start=False)

```
1 print("Random Forest - Test set")
2 rf_score = skl_evaluate(rf, desired_columns, X_test, y_test)
3
4 print("Random Forest - Validation set")
5 rf_score_val = skl_evaluate(rf, desired_columns, X_val, y_val)
```

➞

Random Forest - Test set				
	precision	recall	f1-score	support
0	0.65	0.79	0.72	241
1	0.64	0.46	0.54	188
accuracy			0.65	429
macro avg	0.64	0.63	0.63	429
weighted avg	0.65	0.65	0.64	429

Accuracy on dataset: 0.648019

Feature ranking:

1. Education (0.275820)
2. Age (Birthday Masked) (0.238993)
3. MAX(Learner Test Score) (0.222282)
4. Income (0.123961)
5. Primary Interest In Course (0.076364)
6. How Many Hours A Week Can You Commit To Class (0.034888)
7. Hours Coded (0.027692)

Random Forest - Validation set				
	precision	recall	f1-score	support
0	0.78	0.85	0.82	308
1	0.51	0.39	0.44	121

▼ Random Forest Gridsearch Optimisation

weighted avg	0.70	0.72	0.71	429
--------------	------	------	------	-----

```

1 rf=RandomForestClassifier(random_state=42)
2
3 param_grid = {
4     'n_estimators': [100, 120, 140, 160, 180, 200, 220, 240, 260],
5     'min_samples_leaf': [3, 5, 7],
6     'min_samples_split': [2, 3, 4, 5, 6],
7     'max_depth' : [5, 10, 15, 20, 25, 30, 35, 40, 45],
8     'criterion' :['gini', 'entropy']
9 }
10
11
12 CV_rf = GridSearchCV(estimator=rf, param_grid=param_grid, cv= 5, n_jobs = -1,
13                      verbose = 2)
14 CV_rf.fit(X_train, y_train.astype(int))
15 CV_rf.best_params_

```

This is sample output from this step:

```

Fitting 5 folds for each of 2430 candidates, totalling 12150 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 37 tasks      | elapsed: 8.8s
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed: 33.4s
[Parallel(n_jobs=-1)]: Done 361 tasks     | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 644 tasks     | elapsed: 2.3min
[Parallel(n_jobs=-1)]: Done 1009 tasks    | elapsed: 3.7min
[Parallel(n_jobs=-1)]: Done 1454 tasks    | elapsed: 5.5min
[Parallel(n_jobs=-1)]: Done 1981 tasks    | elapsed: 7.7min
[Parallel(n_jobs=-1)]: Done 2588 tasks    | elapsed: 10.2min
[Parallel(n_jobs=-1)]: Done 3277 tasks    | elapsed: 13.0min
[Parallel(n_jobs=-1)]: Done 4046 tasks    | elapsed: 16.2min

```

```

[Parallel(n_jobs=-1)]: Done 4897 tasks      | elapsed: 19.7min
[Parallel(n_jobs=-1)]: Done 5828 tasks      | elapsed: 23.5min
[Parallel(n_jobs=-1)]: Done 6841 tasks      | elapsed: 27.4min
[Parallel(n_jobs=-1)]: Done 7934 tasks      | elapsed: 32.2min
[Parallel(n_jobs=-1)]: Done 9109 tasks      | elapsed: 37.4min
[Parallel(n_jobs=-1)]: Done 10364 tasks     | elapsed: 43.0min
[Parallel(n_jobs=-1)]: Done 11701 tasks     | elapsed: 49.0min
[Parallel(n_jobs=-1)]: Done 12150 out of 12150 | elapsed: 50.9min finished

```

The optimum parameters are: {'criterion': 'gini', 'max_depth': 5, 'min_samples_leaf': 5, 'min_samples_split': 2, 'n_estimators': 220}

```

1 CV_rf_score = CV_rf.score(X_test, y_test.astype(int))
2 print("Optimised Random Forest accuracy on test:\t %f" % CV_rf_score)

```

▼ Random Forest - Find Optimal Parameters

```

1 def skl_run_combos(combo_columns, model, df_data, df_labels, SHUFFLE, FILL):
2
3     cool_list = list()
4
5     X_train_combo = X_train
6     X_test_combo = X_test
7
8     y_train_combo = y_train
9     y_test_combo = y_test
10
11     for L in range(0, len(combo_columns) + 1):
12         combinations = list(itertools.combinations(combo_columns, L))[1:]
13         for combo in combinations:
14             combo = list(combo)
15
16             model.fit(X_train_combo[combo], y_train)
17             rf_score = model.score(X_test_combo[combo], y_test)
18             output = str(rf_score) + " " + str(combo)
19             print(output)
20             cool_list.append(output)
21
22     cool_list.sort()
23
24     with open('outputs.txt', 'w') as f:
25         for item in cool_list:
26             f.write("%s\n" % item)
27
28
29 def normalise_df(train_df, test_df, AXIS, val_df=1):
30     mu = train_df.mean(axis=AXIS)
31     sd = train_df.std(axis=AXIS)
32
33     train_df = (train_df - mu) / sd
34     test_df = (test_df - mu) / sd
35     val_df = (val_df - mu) / sd
36
37     return train_df, test_df, val_df

```

```

1 df_data = pickle.load( open( "df_data.p", "rb" ) )
2 df_labels = pickle.load( open( "df_labels.p", "rb" ) )
3
4 combo_columns = ['Age (Birthday Masked)', 'Income', 'Education',

```

```

5         'MAX(Learner Test Score)', 'Primary Interest In Course',
6         'Hours Coded', 'How Many Hours A Week Can You Commit To Class',
7         'Promise Zone Indicator', 'Hacker Rank Score',
8         'Location; number']
9
10 model = RandomForestClassifier(criterion='gini', max_depth=5,
11                               min_samples_leaf=5, min_samples_split=2,
12                               n_estimators = 220, oob_score=True,
13                               max_features=0.5, n_jobs = -1, random_state=42)
14
15 print('Accuracy on test, [columns used]')
16 skl_run_combos(combo_columns, model, df_data, df_labels,
17                SHUFFLE, FILL)

```

This is a sample output from this step:

```

0.648018648 'Income', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Promise Zone Indicator', 'Location; number']
0.648018648 'Age (Birthday Masked)', 'Income', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Promise Zone Indicator', 'Location; number']
0.648018648 'Age (Birthday Masked)', 'Income', 'Education', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Promise Zone Indicator', 'Hacker Rank Score', 'Location; number']
0.645687646 'Income', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Promise Zone Indicator', 'Hacker Rank Score', 'Location; number']
0.645687646 'Age (Birthday Masked)', 'Income', 'Education', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Promise Zone Indicator', 'Hacker Rank Score', 'Location; number']
0.643356643 'Income', 'Hours Coded', 'Hacker Rank Score', 'Location; number']
0.643356643 'Age (Birthday Masked)', 'Income', 'Education', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class']
0.643356643 'Age (Birthday Masked)', 'Education', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Hacker Rank Score', 'Location; number']
0.641025641 'Income', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Hacker Rank Score', 'Location; number']
0.641025641 'Age (Birthday Masked)', 'Income', 'Education', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Hacker Rank Score']
0.641025641 'Age (Birthday Masked)', 'Education', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Hacker Rank Score']
0.641025641 'Income', 'Education', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Promise Zone Indicator']
0.641025641 'Income', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Hacker Rank Score', 'Location; number']
0.641025641 'Age (Birthday Masked)', 'Income', 'Education', 'MAX(Learner Test Score)', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Location; number']
0.641025641 'Age (Birthday Masked)', 'Education', 'Hours Coded', 'How Many Hours A Week Can You Commit To Class', 'Promise Zone Indicator', 'Hacker Rank Score', 'Location; number']

```

▼ Multi Layer Perceptron

```

1 def build_mlp(input, NUM_CLASSES):
2     model = keras.Sequential([
3         keras.layers.Dense(8, activation='relu', input_shape=(input.shape[1],)),
4         keras.layers.Dense(5, activation='relu'),
5         # keras.layers.Dense(3, activation='relu'),
6         keras.layers.Dense(NUM_CLASSES, activation='softmax')
7     ])
8
9     model.compile(loss='categorical_crossentropy', optimizer='adam',
10                  metrics=['mae', 'acc'])
11
12     model.summary()
13
14     return model
15
16
17 def train_mlp(dataframe, labels, model, checkpoint_path, EPOCHS):
18     # Limit the training when there are multiple epochs with little change loss
19     # The patience parameter is the amount of epochs to check for improvement
20     early_stop = keras.callbacks.EarlyStopping(monitor='val_loss',
21                                                patience=4000)
22
23     # Create checkpoint callback
24     cp_callback = keras.callbacks.ModelCheckpoint(checkpoint_path,
25                                                  save_weights_only=True,
26                                                  verbose=0)
27
28     # Track the training statistics
29     history = model.fit(dataframe, labels, epochs=EPOCHS,
30                        validation_split=0.2, verbose=0,
31                        callbacks=[early_stop, EpochTrack(), cp_callback])
32
33     print("\nEpochs: {}".format(len(history.epoch)))
34     plot_history(history)
35
36     return model

```

```

1 y_train_sparse = keras.utils.to_categorical(y_train, NUM_CLASSES)

```



```
2
3 TRAIN = 0
4
5 if TRAIN == 1:
6     checkpoint_path = "training_1/cp.ckpt"
7     checkpoint_dir = os.path.dirname(checkpoint_path)
8
9     # Train MLP and save to checkpoints
10    %time mlp = train_mlp(X_train, y_train_sparse, build_mlp(X_train, NUM_CLASSES), cl
11
12 else:
13     checkpoint_path = "training_1/cp.ckpt"
14
15     checkpoint_dir = os.path.dirname(checkpoint_path)
16
17     # Load trained weights from the checkpoint path
18     mlp = build_mlp(X_train, NUM_CLASSES)
19     mlp.load_weights(checkpoint_path)
20
21 # Test
22 print("\nMLP - Test set")
23 acc_test, preds_test = eval_model(mlp, X_test, y_test, 0, NUM_CLASSES)
24
25 print("\nMLP - Validation set")
26 acc_val, preds_val = eval_model(mlp, X_val, y_val, 0, NUM_CLASSES)
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/op
Instructions for updating:
Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
--------------	--------------	---------

▼ MLP every combination of parameters

```

1 col_names = ['Age (Birthday Masked)', 'Income', 'Education',
2             'MAX(Learner Test Score)', 'Primary Interest In Course',
3             'Hours Coded', 'How Many Hours A Week Can You Commit To Class'
4             , 'Promise Zone Indicator', 'Hacker Rank Score', 'Location; number']
5
6 cool_list = list()
7
8 y_train_sparse = keras.utils.to_categorical(y_train, NUM_CLASSES)
9 y_val_sparse = keras.utils.to_categorical(y_val, NUM_CLASSES)
10
11 # with open('results.txt', 'w') as f:
12 #     for L in range(0, len(col_names)+1):
13 #         combinations = list(itertools.combinations(col_names, L))[4:]
14 #         for combo in combinations:
15 #             combo = list(combo)
16
17 #             mlp = train_mlp(X_train[combo], y_train_sparse,
18 #                             build_mlp(X_train[combo], NUM_CLASSES),
19 #                             checkpoint_path, EPOCHS=1000)
20 #             acc_val = eval_model(mlp, X_val[combo], y_val_sparse, 0, NUM_CLASSES)
21
22 #             output = str(acc_val) + " " + str(combo) + "\n"
23 #             print(output)
24 #             cool_list.append(output)
25 #             f.write(str(cool_list))

```

MLP - Validation set

▼ Test Results Summary

```

1 print("Naive Bayes accuracy on validation:\t\t %f" % nb_score_val)
2 print('')
3 print("SVM accuracy on validation:\t\t\t %f" % clf_score_val)
4 print('')
5 print("Linear Regression accuracy on validation:\t %f" % linr_score_val)
6 print('')
7 print("Random Forest accuracy on validation:\t\t %f" % rf_score_val)
8 print('')
9 print("MLP accuracy on validation:\t\t\t %f" % acc_val)

```

Naive Bayes accuracy on validation:	0.641026
SVM accuracy on validation:	0.727273
Linear Regression accuracy on validation:	0.051558
Random Forest accuracy on validation:	0.722611
MLP accuracy on validation:	0.668998

▼ Insight

Let's sanity check the results by comparing model predictions to the real test results.

```

1 predictions = preds_test
2
3 y_classes = predictions.argmax(axis=-1)

```

When comparing predictions to true labels we'd expect to see mostly true values for high scoring classifiers

```

1 y_classes = y_classes.flatten()
2
3 # Match is 'True' when values are the same
4 match = [y_classes[x] == y_test[x] for x in range(len(y_classes))]
5
6 # Combine and show
7 df_compare = pd.DataFrame({'Predictions': y_classes, 'Actual': y_test,
8                           'Match': match})
9
10 print('Show performance on first 5 rows')
11 df_compare.head(5)

```

☞ Show performance on first 5 rows

	Predictions	Actual	Match
0	1	1	True
1	0	0	True
2	1	1	True
3	0	1	False
4	1	1	True

We can also obtain the set of rows that were incorrectly predicted

```

1 df_pred = show_incorrect(y_classes, X_test, y_test)
2
3 df_pred.sort_values('Prediction', inplace=True)
4 print('Show first 5 rows that had wrong predictions')
5 print('Values are encoded as per earlier steps')
6 df_pred.head(5)

```

☞ Show first 5 rows that had wrong predictions
Values are encoded as per earlier steps

	Age (Birthday Masked)	Income	MAX(Learner Test Score)	Hours Coded	How Many Hours A Week Can You Commit To Class	Promise Zone Indicator	Location; number	p
0	1.647899	1.832671	-0.628481	-0.055229	0.381102	-0.405866	0.571825	
125	-0.185411	-1.578677	0.988618	-0.055229	-0.823068	-0.405866	0.571825	
124	-0.293252	-0.299421	0.341779	1.974047	0.381102	-0.405866	0.571825	

